

# Technical Overview: Idea Validator

## Business Content & Problem statement

### What problem does my agent solve?

The problem my agent solves is validating ideas, whether that be for business, apps, startups and more. I have personally dealt with this issue, as anytime I have come up with an idea I would have to make multiple searches and even then would sometimes miss a key competitor. Even worse, there may have been situations where I have missed a key detail that could make my idea that already had exposure unique. This aspect is important as many of the best products aren't even brand new, however they are an existing system with one or two key changes, or combinations, to vastly improve the experience.

### Why does it benefit from AI?

Idea validation benefits greatly from the use of AI. People with new ideas can now quickly access information about competitors, existing market niches, and potential strategies for differentiation. The agent acts like a strategic assistant, offering analytical feedback and even suggesting roadmaps for how to begin developing the idea, a critical support during the often isolated early phase of launching an idea into the world.

## Architecture Justification

### Frontend Architecture & UX Design Rationale

The frontend architecture is component based, making sure to separate concerns. Furthermore, I opted for a SPA rather than implementing navigation, as the current experience can fill one page very suitably. This keeps the UI focused and fast. The UI takes inspiration from existing LLM interfaces like OpenAI Chat and Claude, designing the chat-style interface to offer users a familiar and intuitive experience, tailored specifically for idea innovation and validation. Rather than reinventing the wheel, I focused on building upon established UX patterns to ensure users feel comfortable and confident using the tool from the start, which is especially important for engagement and retention. Finally, I implemented full support for both dark and light modes to accommodate different user visual preferences and enhance accessibility.

### LLM Integration, Model selection & Prompting

To integrate the LLM, I built a Flask backend with a single route for chat submissions. This endpoint receives the current conversation history and new user message, handles moderation checks, and streams responses using a custom queue-based handler. It also supports tool calling and dynamic context trimming to fit within reasonable pricing limits, providing a fast and responsive chat experience. I chose GPT-4o as the model due to its combination of high performance, large context window for memory retention, and direct support for streaming responses. Its compatibility with LangChain agent system made it ideal for building a tool-augmented, highly responsive application. For prompting, my agent takes a structured approach. A system prompt is prepended to every interaction, clearly instructing the model on tool calling behavior, safety guidelines, when to ask follow up questions, and more. This is followed by the first user-assistant message pair to anchor the session, along with a trimmed portion of recent conversation history. This keeps context relevant while controlling token usage and cost. The user's latest message is appended at the end. This design ensures consistent behavior and reinforces the agent's specialized role.

### Tool Integration & Agent Design

For tool integration, I went with Tavily Search to allow the agent to perform live web lookups to support business idea validation, and a custom "Generate Roadmap" tool that provides the user with a key-point summarization of their idea and where to go next. The tools are passed to the agent through the `initialize_agent()` function using OpenAI's Functions agent type, enabling the LLM to decide when to invoke them based on context and prompting. Regarding the agent design, I opted for a setup that delegates decision making to the model. This includes tool usage, allowing the model to determine when a

live web lookup or roadmap generation are necessary, making the experience feel natural. The agent also includes a content moderation layer using OpenAI's Moderation endpoint to block harmful or policy violating input before processing.

## **Technical Design Decisions**

### **Performance Considerations for Enterprise Development**

To ensure strong performance under enterprise level usage, several optimizations were implemented across both the frontend and backend. On the backend, token usage is controlled by dynamically trimming the conversation context to stay within budget, and responses are streamed to improve perceived latency and responsiveness. On the frontend, visual cues are shown during tool invocation to set user expectations for additional processing time. React's `useRef()` is used to manage DOM updates efficiently, avoiding unnecessary re-renders.

### **Error Handling & Reliability Measures**

For error handling, the app surfaces them directly in the agent's response stream, following a pattern commonly used in production-grade AI interfaces. I accounted for several types of failures, including connection issues, agent-side errors, timeouts, and content violation. Each type of error has its own short, descriptive message to inform the user without overly disrupting the experience. To improve reliability, I implemented guardrails in the UI such as a maximum text input length, disabling the submit button while the agent is responding, and enforcing a 60 second timeout on responses. These choices help ensure consistency and prevent overload, aligning with production grade frontend development practices.

The current implementation is optimized for small scale testing but would require several upgrades for production scalability. At the moment, conversation history is passed directly in each request without persistent storage in the backend. In a real world deployment, this should be backed by a database rather than local memory to support concurrent users and long-term state retention. Token usage is another key concern. While GPT-4o is relatively affordable for development, production-scale usage could lead to high API costs, especially given the model's large context window. Cost control strategies such as message summarization or caching common search results would help manage this.

## **Innovation and Creativity**

### **Novel Application of AI**

Most AI applications follow a general purpose, "ask me anything" model, which often leads to unfocused or surface-level interactions. This application stands out by offering a tightly focused experience centered on business idea validation. Its purpose driven design ensures that every response is anchored in helping users refine and develop their ideas. Furthermore, the agent dynamically decides when to use an integrated web search tool, allowing it to pull in relevant, real time information when needed. The result is a lightweight yet intelligent assistant that delivers targeted insights.

### **Future enhancement ideas**

As an engineer, I am always thinking about how I can further build to improve the experience for my users. This stays true with this endeavor as many ideas have and continue to circulate:

1. Customizable experience section with options on how to curate their assistant, with differing system prompts for each.
2. A function that auto generates a title for the user's idea based on conversation context, rather than the first 30 characters.
3. Let users export their roadmaps as a PDF.

## **Documentation of AI use**

I used AI tools to help with various aspects of development. Areas where it helped me the most were drafting the system prompt, evaluating tradeoffs in model and tool integration, and designing the token streaming logic like using the custom queue. In addition, it helped me discover unfamiliar concepts, like React's version of directly injecting HTML, allowing me to stylize the text of my agent's output.